

# Predictive Models of the Age of Abalones Based on Physical Characteristics

## Summary

In this report, we aim to create an effective model which can predict age of Abalone accurately based on individual physical characteristics. We perform exploratory data analysis and visualization, test both linear and non-linear models on our data, then compare models through regression metrics.

Authors: Serene Zha, Mehmet Imgä, Claudia Liauw, Wendy Frankel

## Introduction

Abalone are marine mollusks that are commercially important in fisheries and aquaculture, particularly in regions such as Tasmania. Estimating the age structure of abalone populations is essential for setting sustainable harvest limits and monitoring stock health. However, the standard method for determining age requires cutting the shell through the cone, staining it, and counting growth rings under a microscope—a destructive, time-consuming, and labor-intensive procedure[1]. Because of this, methods that can infer age from simple, non-destructive measurements of the animal are of practical interest to biologists, fisheries managers, and growers.

Here, we ask whether we can use a machine learning model to predict the age of an abalone from basic physical measurements. Specifically, we will explore linear regression models in Python to relate age to various attributes, including sex, shell length, diameter, height, and several weight measurements. To investigate this question, we use the UCI Abalone dataset, which contains 4,177 abalones with eight predictor variables and a target variable, "Rings." Each row corresponds to one abalone, and the recorded features include sex (male, female, infant), three shell size measurements (length, diameter, height), and four weight measurements (whole, shucked, viscera, and shell weight). The number of rings serves as a proxy for age, with age in years given approximately by  $\text{Rings} + 1.5$  [2], to account for any rings missed in counting (they can be difficult to see). By building and evaluating linear regression models on this dataset, we aim to understand how well these readily obtained physical measurements can predict abalone age and what this implies for practical, non-destructive age estimation.

## Methods and Results

## 1. Load Necessary Packages

```
In [1]: from ucimlrepo import fetch_ucirepo
import pandas as pd
import altair as alt
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.compose import make_column_transformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler, PolynomialFeatures
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score
```

## 2. Load Data

```
In [2]: import requests
import zipfile

url = "https://archive.ics.uci.edu/static/public/1/abalone.zip"

request = requests.get(url)
with open("../data/raw/abalone.zip", 'wb') as f:
    f.write(request.content)

with zipfile.ZipFile("../data/raw/abalone.zip", 'r') as zip_ref:
    zip_ref.extractall("../data/raw")
```

```
In [3]: # fetch dataset
abalone = fetch_ucirepo(id=1)

# Extract features and targets
X = abalone.data.features
y = abalone.data.targets
```

## 3. Data Wrangling and Cleaning

### Split Data

```
In [4]: # Split Data (Same random_state as baseline for comparison)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=522
)

train_df = pd.concat([X_train, y_train], axis=1)
test_df = pd.concat([X_test, y_test], axis=1)

# ravel y for sklearn
y_train = y_train.values.ravel()
y_test = y_test.values.ravel()
```

```
print("Data loaded and split.")
```

```
train_df
```

Data loaded and split.

Out[4]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	She
2194	I	0.430	0.325	0.110	0.3675	0.1355	0.0935	
3996	I	0.315	0.230	0.000	0.1340	0.0575	0.0285	
3329	F	0.545	0.435	0.150	0.6855	0.2905	0.1450	
492	F	0.655	0.510	0.155	1.2895	0.5345	0.2855	
241	I	0.270	0.200	0.070	0.1000	0.0340	0.0245	
...	...	...	...	...	...	...	...	...
3956	F	0.515	0.395	0.140	0.6860	0.2810	0.1255	
154	F	0.565	0.450	0.135	0.9885	0.3870	0.1495	
3360	F	0.580	0.440	0.175	1.0730	0.4005	0.2345	
1899	M	0.575	0.450	0.130	0.7850	0.3180	0.1930	
3988	M	0.665	0.515	0.165	1.3855	0.6210	0.3020	

3341 rows × 9 columns

## Save Data

In [5]:

```
train_df.to_csv('../data/processed/abalone_train.csv', index=False)  
test_df.to_csv('../data/processed/abalone_test.csv', index=False)
```

## 4. EDA Summary

- Mostly numerical variables except sex.
- No missing values.
- Target (Rings) ranges from 1 to 29. Mostly normal, slight right skew.
- Sex needs to be one-hot encoded, the rest should be scaled.
- Numeric variables are moderately positively correlated with target.

In [6]:

```
missing_values = train_df.isnull().sum()  
print("Missing values per column:", missing_values[missing_values > 0])
```

Missing values per column: Series([], dtype: int64)

In [7]:

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3341 entries, 2194 to 3988
Data columns (total 9 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Sex               3341 non-null    object  
 1   Length            3341 non-null    float64 
 2   Diameter          3341 non-null    float64 
 3   Height            3341 non-null    float64 
 4   Whole_weight      3341 non-null    float64 
 5   Shucked_weight   3341 non-null    float64 
 6   Viscera_weight   3341 non-null    float64 
 7   Shell_weight     3341 non-null    float64 
 8   Rings             3341 non-null    int64  
dtypes: float64(7), int64(1), object(1)
memory usage: 261.0+ KB
```

```
In [8]: train_df.describe().round(2)
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_w
<b>count</b>	3341.00	3341.00	3341.00	3341.00	3341.00	3341.00	3341.00
<b>mean</b>	0.52	0.41	0.14	0.82	0.36	0.18	0.76
<b>std</b>	0.12	0.10	0.04	0.49	0.22	0.11	0.35
<b>min</b>	0.08	0.06	0.00	0.00	0.00	0.00	0.00
<b>25%</b>	0.45	0.35	0.12	0.44	0.18	0.09	0.58
<b>50%</b>	0.55	0.42	0.14	0.80	0.33	0.17	0.76
<b>75%</b>	0.62	0.48	0.16	1.14	0.50	0.25	0.94
<b>max</b>	0.82	0.65	0.52	2.83	1.49	0.76	1.00

## 5. EDA Visualisation

```
In [9]: from ydata_profiling import ProfileReport
ProfileReport(train_df)
```

[Upgrade to ydata-sdk](#)

Improve your data and profiling with ydata-sdk, featuring data quality scoring, redundancy detection, outlier identification, text validation, and synthetic data generation.

Summarize dataset: 0% | 0/5 [00:00<?, ?it/s]

100% | ██████████ | 9/9 [00:00<00:00, 791.30it/s]

Generate report structure: 0% | 0/1 [00:00<?, ?it/s]

Render HTML: 0% | 0/1 [00:00<?, ?it/s]

# Overview

Brought to you by [YData](#)

[Overview](#)[Alerts 8](#)[Reproduction](#)

## Dataset statistics

**Number of variables**

9

**Number of observations**

3341

**Missing cells**

0

**Missing cells (%)**

0.0%

**Duplicate rows**

0

**Duplicate rows (%)**

0.0%

**Total size in memory**

261.0 KiB

**Average record size in memory**

80.0 B

## Variable types

**Categorical**

1

**Numeric**

8

# Variables

Out[9]:

## Further EDA Visualisation

In [10]:

```
# Correlation matrix heatmap
corr_matrix = train_df.select_dtypes(include=['float64', 'int64']).corr().reset_index
corr_df = pd.melt(corr_matrix, id_vars='index', var_name='variable2', value_name='c
```

```

heatmap = alt.Chart(corr_df).mark_rect().encode(
    x=alt.X('index', title=None),
    y=alt.Y('variable2', title=None),
    color=alt.Color('correlation', scale=alt.Scale(scheme='blueorange', domain=[-1,
        1]),
    tooltip=['index', 'variable2', 'correlation']
).properties(
    title='Correlation Matrix',
    width=400,
    height=400
)

heatmap

```

Out[10]:

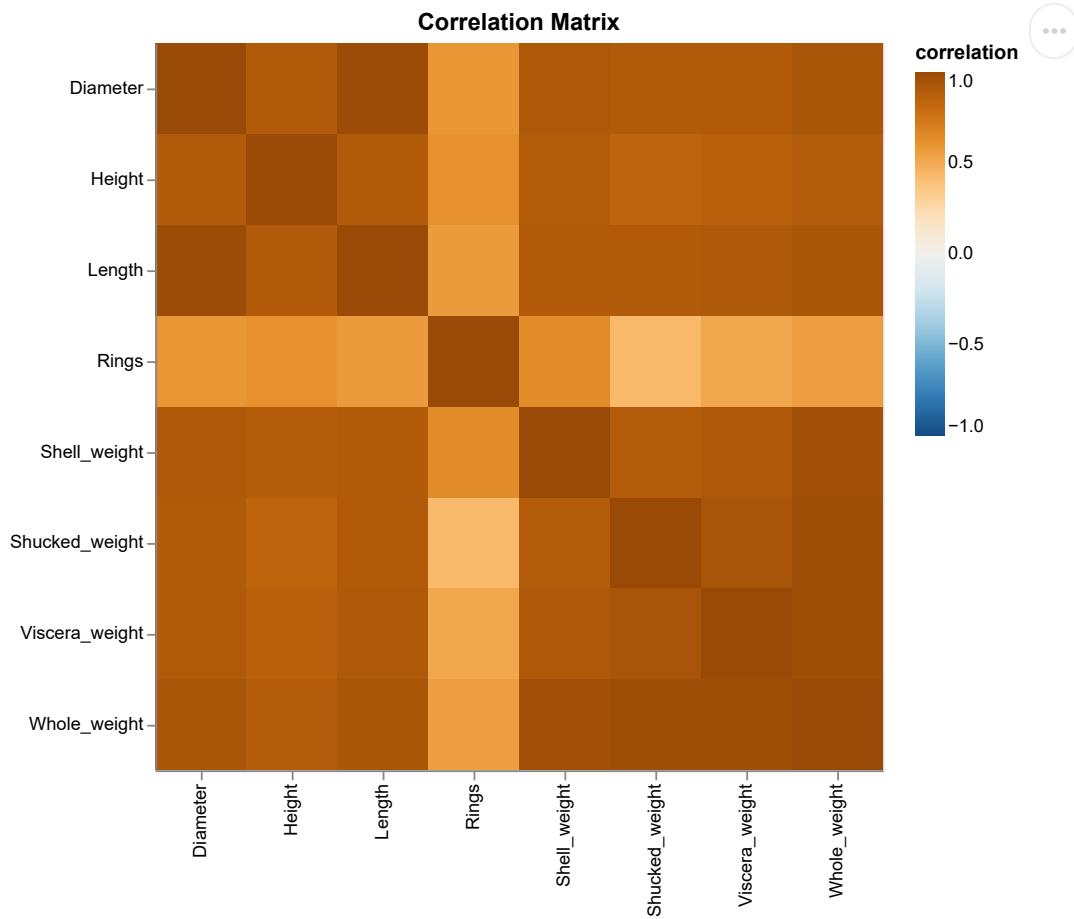


Figure 1: Heatmap showing the correlation between different numerical features and the target variable Rings.

Below, we investigate the possible relationship between sex of adults (M/F), Infants, and number of rings, as the relationship may differ between those categories.

```

In [11]: base = alt.Chart(train_df).mark_circle(opacity=0.3).encode(
    x='Shell_weight',
    y='Rings',
    color='Sex'
)

lines = base.transform_regression(

```

```

    'Shell_weight', 'Rings', groupby=['Sex']
).mark_line().encode(
  color='Sex'
)

(base + lines).properties(
  title="Rings vs Shell Weight by Sex (with Regression Lines)",
  width=500
)
)

```

Out[11]:

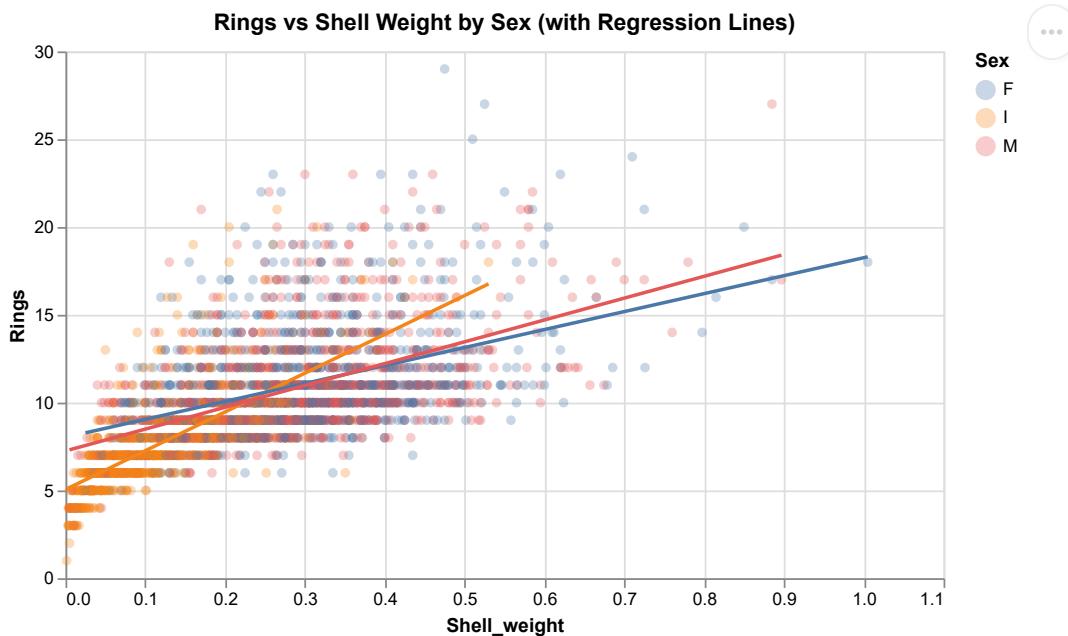


Figure 2. Number of rings (target variable) by shell weight (grams), grouped into Male (M), Female (F), and Infant (I). Linear regression lines for each group are shown.

Observation: The slopes appear slightly different, particularly for Infants (I) compared to adults (M/F). Infants seem to have a steeper growth curve in this dimension.

## 6. Modeling Interactions

We will use a Linear Regression model to predict the number of rings. We construct a pipeline that:

1. One-hot encodes the categorical Sex feature.
2. Scales the numerical features.
3. Applies Linear Regression.

In [12]:

```

# Define features
numeric_features = ['Length', 'Diameter', 'Height', 'Whole_weight', 'Shucked_weight'
categorical_features = ['Sex']

# Create preprocessor
preprocessor = make_column_transformer(
  (StandardScaler(), numeric_features),
  (OneHotEncoder(drop='if_binary'), categorical_features)
)

```

```

)

# Create and fit pipeline
lr_model = make_pipeline(preprocessor, LinearRegression())
lr_model.fit(X_train, y_train)

# Make predictions
lr_y_pred = lr_model.predict(X_test)

```

## 7. Visual Evaluation of the model:

We evaluate the model by plotting the Predicted vs. Actual values.

```
In [13]: # Calculate metrics
lr_rmse = np.sqrt(mean_squared_error(y_test, lr_y_pred))
lr_r2 = r2_score(y_test, lr_y_pred)

print(f"Root Mean Squared Error (RMSE): {lr_rmse:.4f}")
print(f"R-squared (R2): {lr_r2:.4f}")

# Visualization
results_df = pd.DataFrame({
    'Actual': y_test.flatten(),
    'Predicted': lr_y_pred.flatten()
})

pred_chart = alt.Chart(results_df).mark_circle(opacity=0.5).encode(
    x=alt.X('Actual', title='Actual Rings'),
    y=alt.Y('Predicted', title='Predicted Rings')
).properties(
    title=f'Actual vs Predicted Rings (R2 = {lr_r2:.2f})',
    width=500,
    height=500
)

line = alt.Chart(pd.DataFrame({'x': [0, 30], 'y': [0, 30]})).mark_line(color='red',
    x='x',
    y='y'
)

pred_chart + line
```

Root Mean Squared Error (RMSE): 2.3419  
R-squared (R2): 0.4427

Out[13]:

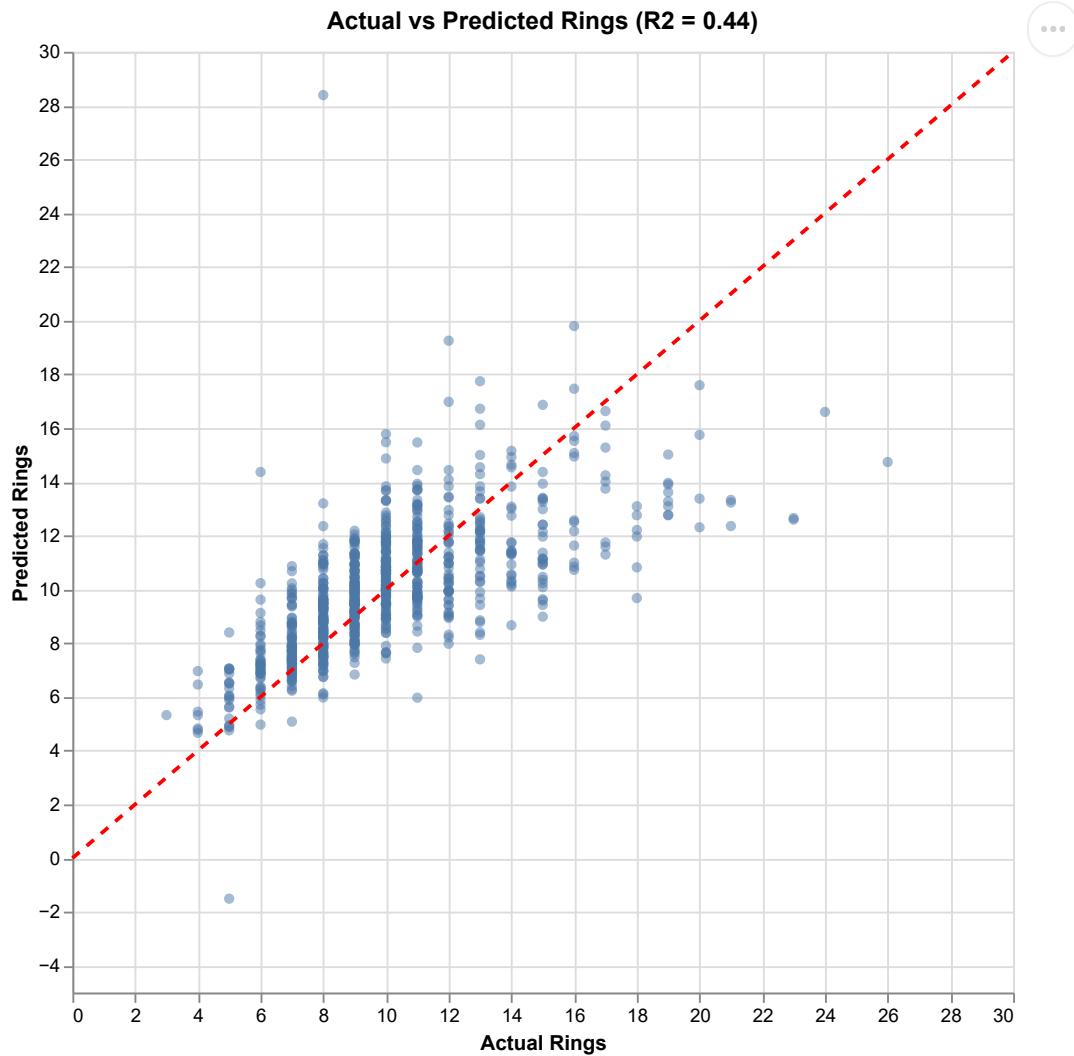


Figure 3: Actual Rings (x-axis) vs Predicted Rings (y-axis). The red dashed line represents the linear regression line. Points below the line indicate over-prediction, while points above indicate under-prediction.

In [14]:

```
# Extracting feature names from the preprocessor
ohe = lr_model.named_steps['columntransformer'].named_transformers_['onehotencoder']
ohe_features = list(ohe.get_feature_names_out(categorical_features))

# combining to get all features
all_features = numeric_features + ohe_features

# Extract coefficients from Linear regression
coefficients = lr_model.named_steps['linearregression'].coef_

coef_df = pd.DataFrame({
    "Feature": all_features,
    "Coefficient": coefficients
}).sort_values("Coefficient", ascending=False)

coef_df
```

Out[14]:

	Feature	Coefficient
3	Whole_weight	4.483371
1	Diameter	1.153338
6	Shell_weight	1.020152
2	Height	0.808758
9	Sex_M	0.306958
7	Sex_F	0.208326
0	Length	-0.265891
8	Sex_I	-0.515284
5	Viscera_weight	-1.184490
4	Shucked_weight	-4.397524

Table 1. Linear Regression coefficient values showing the estimated contribution of each feature to the predicted target after preprocessing and one-hot encoding.

## 8. Analysis 2: Non-Linear Models

### Model A: Random Forest Regressor

```
In [15]: rf_preprocessor = make_column_transformer(  
    (StandardScaler(), numeric_features), # Scaling isn't strictly necessary for RF  
    (OneHotEncoder(drop='if_binary'), categorical_features)  
)  
  
rf_model = make_pipeline(rf_preprocessor, RandomForestRegressor(n_estimators=100, r  
rf_model.fit(X_train, y_train)  
y_pred_rf = rf_model.predict(X_test)  
  
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))  
r2_rf = r2_score(y_test, y_pred_rf)  
  
print(f"Random Forest - RMSE: {rmse_rf:.4f}, R2: {r2_rf:.4f}")
```

Random Forest - RMSE: 2.1817, R2: 0.5163

### Model B: Support Vector Regression (SVR)

```
In [16]: svr_preprocessor = make_column_transformer(  
    (StandardScaler(), numeric_features), # Scaling is CRITICAL for SVR  
    (OneHotEncoder(drop='if_binary'), categorical_features)  
)  
  
svr_model = make_pipeline(svr_preprocessor, SVR(kernel='rbf', C=1.0, epsilon=0.1))  
svr_model.fit(X_train, y_train)  
y_pred_svr = svr_model.predict(X_test)
```

```

rmse_svr = np.sqrt(mean_squared_error(y_test, y_pred_svr))
r2_svr = r2_score(y_test, y_pred_svr)

print(f"SVR (RBF Kernel) - RMSE: {rmse_svr:.4f}, R2: {r2_svr:.4f}")

```

SVR (RBF Kernel) - RMSE: 2.1673, R2: 0.5227

## 9. Comparison of Models

```

In [17]: results = pd.DataFrame({
    'Model': ['Baseline (Linear)', 'Random Forest', 'SVR (RBF)'],
    'RMSE': [lr_rmse, rmse_rf, rmse_svr],
    'R2 Score': [lr_r2, r2_rf, r2_svr]
})

print(results.round(4))

# Visualize Comparison
base = alt.Chart(results).encode(x=alt.X('Model', sort='y'))

bar_r2 = base.mark_bar().encode(
    y=alt.Y('R2 Score', title='R2 Score'),
    color=alt.Color('Model', legend=None)
).properties(title='Model Performance Comparison (R2)')

bar_rmse = (
    alt.Chart(results)
    .encode(
        x=alt.X('Model', sort='y'),
        y=alt.Y('RMSE', title='RMSE'),
        color=alt.Color('Model', legend=None)
    )
    .mark_bar()
    .properties(title='Model Performance Comparison (RMSE)')
)

bar_rmse | bar_r2

```

	Model	RMSE	R2 Score
0	Baseline (Linear)	2.3419	0.4427
1	Random Forest	2.1817	0.5163
2	SVR (RBF)	2.1673	0.5227

Out[17]: Model Performance Comparison (RMSE) Model Performance Comparison (R2) 

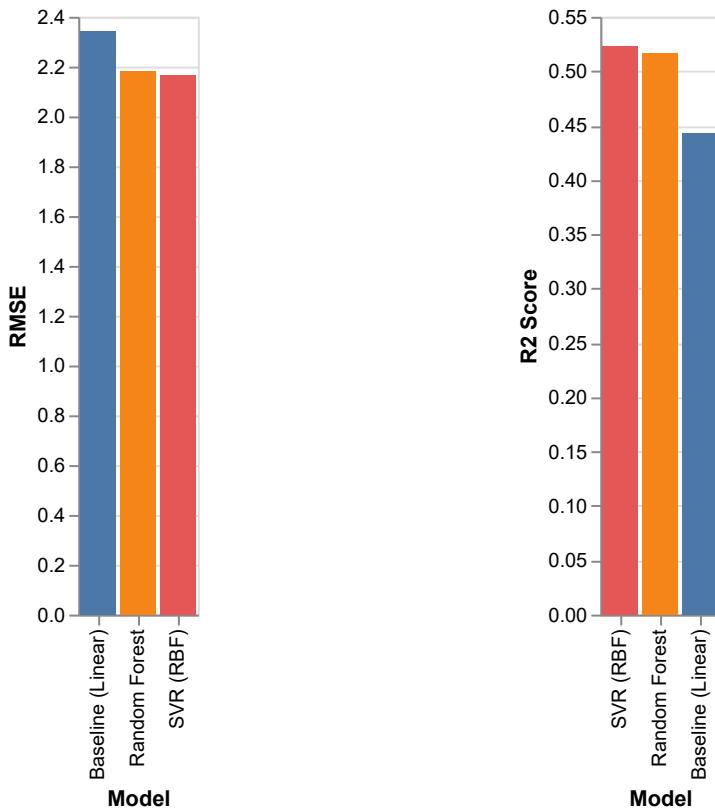


Figure 4. RMSE and  $R^2$  of each model tested.

## Discussion

In this report, we tested 3 different models to find which model was best at predicting number of rings (as a proxy for age) of Abalone molluscs. The baseline linear model explains about half of the variance in number of rings ( $R^2 = 0.44$ ) using size and weight; errors are moderate and roughly symmetric, with a root mean squared error (RMSE) of approximately 5.48. The other two models tested were non-linear models. The Random Forest model generally achieves higher  $R^2$  (0.52) and lower RMSE (4.76) than linear models, showing that abalone growth is not purely linear in the features. Finally, the support vector regression with an RBF kernel model is also an improvement over the baseline (RMSE of 4.70 and  $R^2$  of 0.52), especially in mid-age ranges, but can be more sensitive to scaling and hyperparameters than Random Forest. The non-linear models (especially random forest) provided the best predictive performance, suggesting that future work should focus on flexible models and possibly add environmental features (e.g., location, temperature).

Although the linear model did not perform well relative to the non-linear models, extracting the coefficients for each feature reveals important patterns in our data. Whole weight had the strongest positive coefficient at (4.48), while shucked weight had the lowest coefficient (-4.40). These data indicate two things: that larger, heavier abalone (higher whole weight) had a strong positive relationship with age, and that there is likely collinearity in these

features. We can assume possibility of collinearity based on the fact that shucked weight has a strong negative coefficient, despite the fact that, in context, shucked weight and whole weight should be highly associated. In the future, models which account for this collinearity, and potentially others, should be utilized. That whole weight is a strong positive predictor of age is expected - these organisms continue to grow larger with age throughout their lives [2].

Within the context of biological organisms, that a linear regression model did not perform well is unsurprising. Growth rate, which impacts features such as those in our dataset, can be highly variable depending on biotic and abiotic factors, and so it is likely that a linear regression model cannot account for enough of the variance present in the dataset. In a literature review by Guney et al. (2022), the authors review six different predictive models of Abalone age based on physical characteristics, and found that the BPFFNN model had the highest test accuracy, followed by a random forest model [3]. This agrees with our results on which model performed best.

Estimating Abalone age in a non-destructive way is important for growers and consumers of this mollusc. Therefore, a model which can predict age using easy-to-obtain physical measurements would be of practical importance. Our results suggest that the model should be non-linear, and that weight is likely to be an important predictor of age. However, there are still many improvements to be made on models predicting Abalone age, and questions to be answered. Future work could improve on models by gathering information on other physical characteristics, investigating collinearity in features, and testing other types of models.

## References

- [1] Dua, D., & Graff, C. (2019). UCI Machine Learning Repository: Abalone Data Set. University of California, Irvine, School of Information and Computer Science. Retrieved from the UCI Machine Learning Repository.
- [2] Nash, W. J., Sellers, T. L., Talbot, S. R., Cawthorn, A. J., & Ford, W. B. (1994). The population biology of abalone (*Haliotis* species) in Tasmania. I. Blacklip abalone (*H. rubra*) from the north coast and islands of Bass Strait. Sea Fisheries Division Technical Report No. 48.
- [3] Guney, S., Kilinc, I., Hameed, A.A., Jamil, A. (2022). Abalone Age Prediction Using Machine Learning. In: Djeddi, C., Siddiqi, I., Jamil, A., Ali Hameed, A., Kucuk, İ. (eds) Pattern Recognition and Artificial Intelligence. MedPRAI 2021. Communications in Computer and Information Science, vol 1543. Springer, Cham. [https://doi.org/10.1007/978-3-031-04112-9\\_25](https://doi.org/10.1007/978-3-031-04112-9_25)
- [4] Python Core Team. Python: A dynamic, open source programming language. Python Software Foundation, 2019. Python version 2.7. URL: <https://www.python.org/>.

